Introduction to Azure OpenAI Service

Unlocking the Power of Language Models with Microsoft Azure

Date: April 7, 2025

Ivan Portilla

STAT 5350/4350

What is Azure OpenAI Service?

- REST API access to OpenAI models via Azure
- Supports GPT-4o, GPT-4 Turbo, GPT-3.5, DALL-E, Whisper, etc.
- Use cases:
  - Text & image generation
  - Summarization
  - Code generation
  - Semantic search
  - Image understanding

https://learn.microsoft.com/en-us/azure/ai-services/openai/overview

Key Features Overview

- Models: GPT-4o, GPT-4 T, GPT-3.5, Embeddings
- Vision support: GPT-4 Turbo with Vision
- Fine-tuning: GPT-4, GPT-3.5-Turbo (preview)
- UI: Azure Portal, AI Foundry
- Content filtering: Built-in

Responsible AI with Microsoft

- Microsoft applies Responsible AI principles:
  - Fairness, Safety, Privacy, Transparency
- Tools:
  - Content filters
  - Responsible AI documentation
  - Code of Conduct
  - Limited Access for sensitive features

How to Get Started

1. Create Azure OpenAI Resource
2. Deploy a Model (e.g., GPT-4o)
3. Use the model:
   - AI Foundry Playground
   - REST API or SDKs (Python, C#, JS, etc.)

Lab 1 – Set up Az OpenAI service

https://learn.microsoft.com/en-us/azure/ai-services/openai/how-to/create-resource

1. Sign in to Portal

https://portal.azure.com/

1. Create a resource in Azure Services
2. Search Azure OpenAI

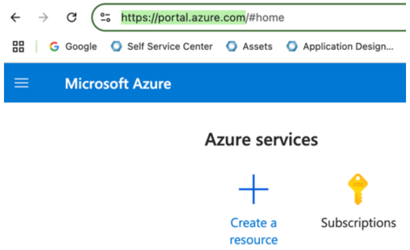

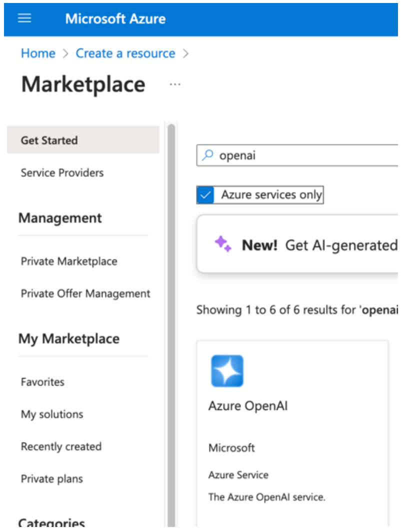

openai

☑ Azure services only

New! Get AI-generated

Showing 1 to 6 of 6 results for 'openai

Azure OpenAI
Microsoft
Azure Service
The Azure OpenAI service.

Lab 1 – Set up Az OpenAI service

3. Complete Basics tab (take defaults and click create)

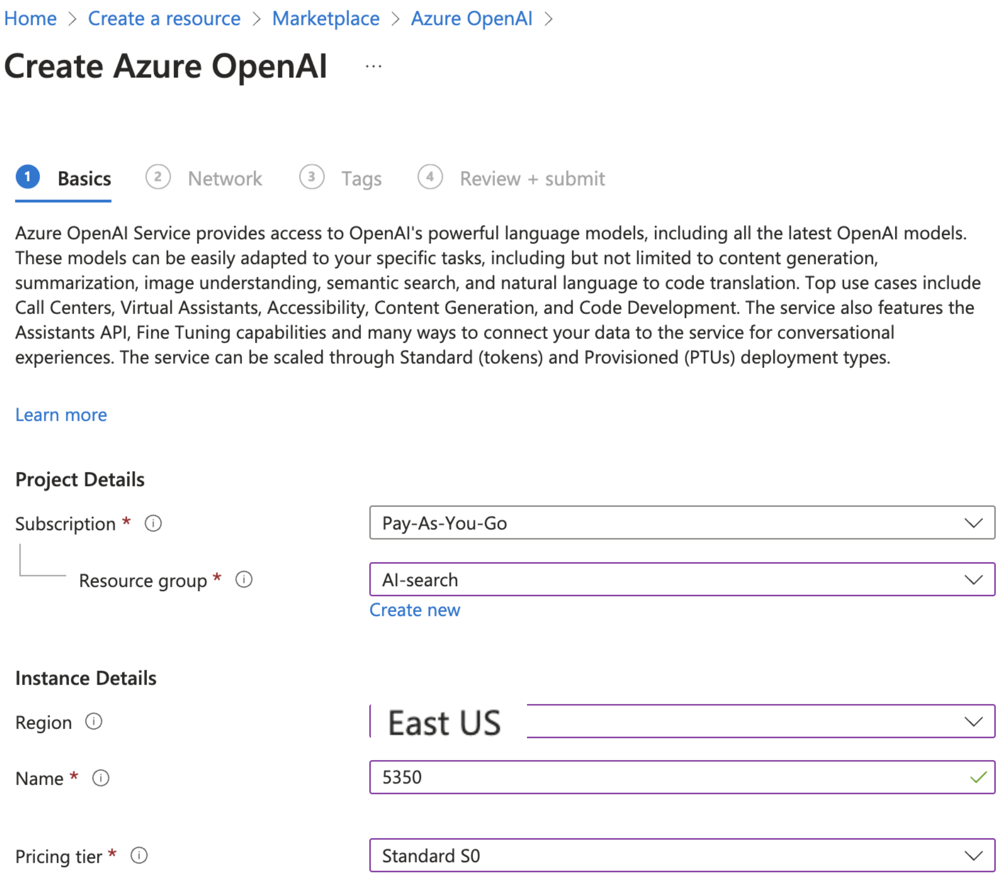

Create Azure OpenAI

Azure OpenAI Service provides access to OpenAI's powerful language models, including all the latest OpenAI models. These models can be easily adapted to your specific tasks, including but not limited to content generation, summarization, image understanding, semantic search, and natural language to code translation. Top use cases include Call Centers, Virtual Assistants, Accessibility, Content Generation, and Code Development. The service also features the Assistants API, Fine Tuning capabilities and many ways to connect your data to the service for conversational experiences. The service can be scaled through Standard (tokens) and Provisioned (PTUs) deployment types.

Learn more

Project Details

Subscription * Pay-As-You-Go

Resource group * AI-search

Create new

Instance Details

Region East US

Name * 5350

Pricing tier * Standard S0

Lab 2 – Chat playground

- Open AI OpenAI Resource (ai.azure.com)
- In Chat Playground, create Deployment

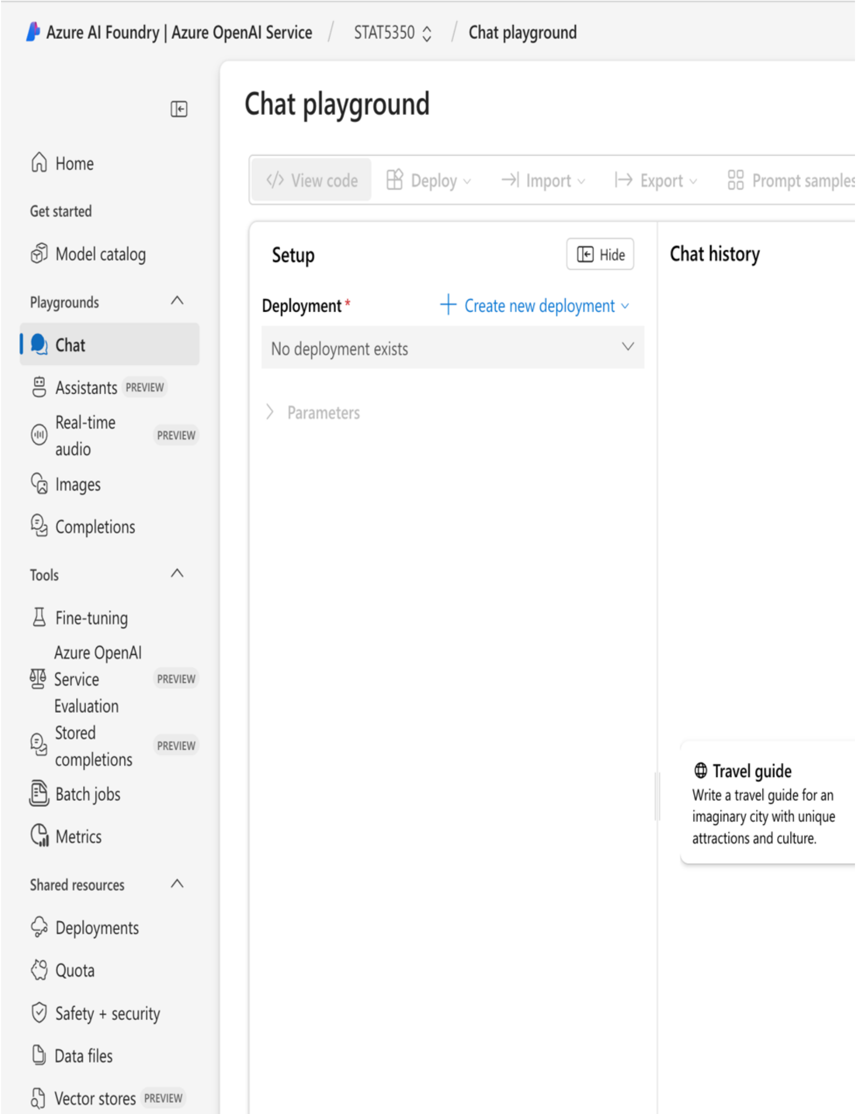

Deployment * No deployment exists

Write a travel guide for an imaginary city with unique attractions and culture.

Lab 2 – Chat playground

- Select gpt-4o-mini
- Click Confirm
- Accept default values
- Click Deploy

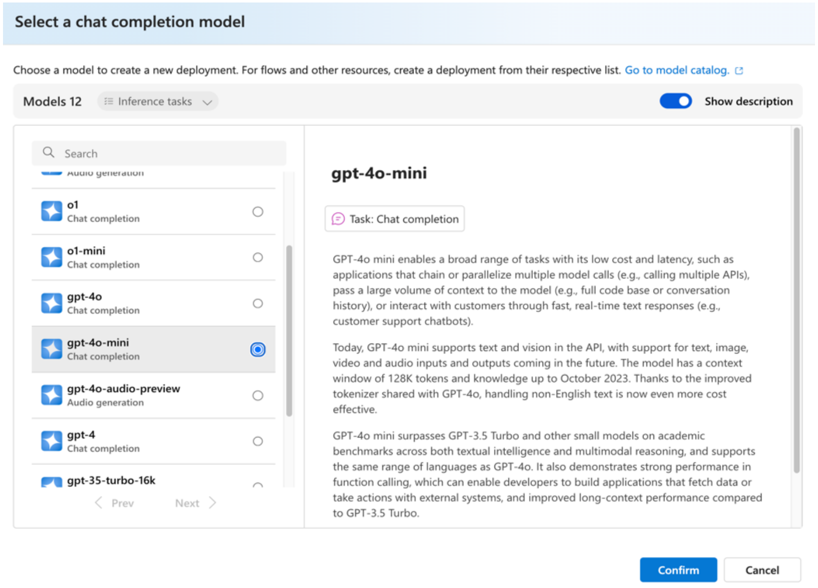

Select a chat completion model

Choose a model to create a new deployment. For flows and other resources, create a deployment from their respective list. Go to model catalog.

Models 12 | Inference tasks | Show description

Search

- Audio generation
- o1 – Chat completion ○
- o1-mini – Chat completion ○
- gpt-4o – Chat completion ○
- gpt-4o-mini – Chat completion ●
- gpt-4o-audio-preview – Audio generation ○
- gpt-4 – Chat completion ○
- gpt-35-turbo-16k

< Prev Next >

gpt-4o-mini

Task: Chat completion

GPT-4o mini enables a broad range of tasks with its low cost and latency, such as applications that chain or parallelize multiple model calls (e.g., calling multiple APIs), pass a large volume of context to the model (e.g., full code base or conversation history), or interact with customers through fast, real-time text responses (e.g., customer support chatbots).

Today, GPT-4o mini supports text and vision in the API, with support for text, image, video and audio inputs and outputs coming in the future. The model has a context window of 128K tokens and knowledge up to October 2023. Thanks to the improved tokenizer shared with GPT-4o, handling non-English text is now even more cost effective.

GPT-4o mini surpasses GPT-3.5 Turbo and other small models on academic benchmarks across both textual intelligence and multimodal reasoning, and supports the same range of languages as GPT-4o. It also demonstrates strong performance in function calling, which can enable developers to build applications that fetch data or take actions with external systems, and improved long-context performance compared to GPT-3.5 Turbo.

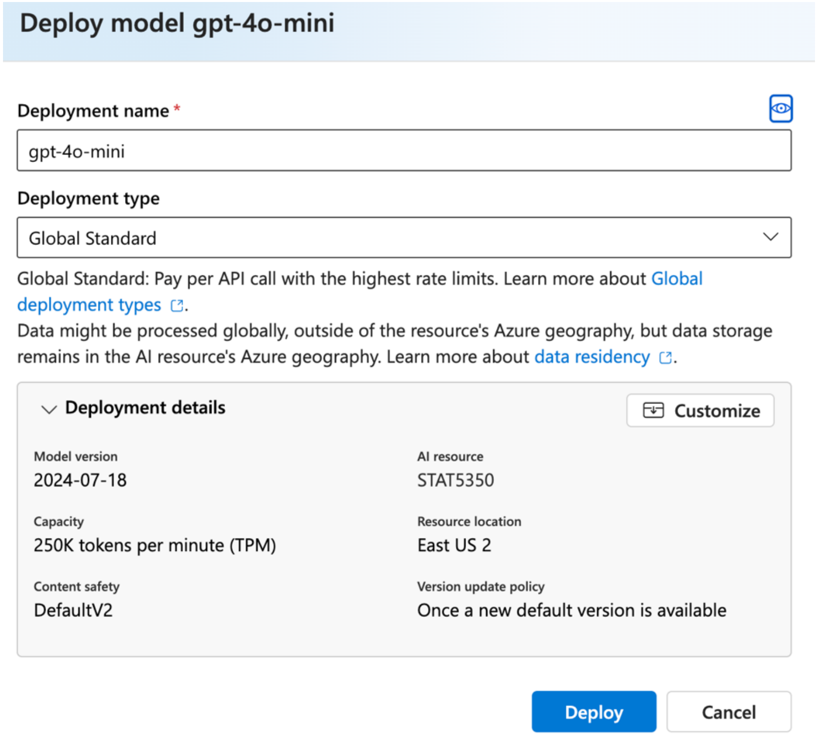

Deploy model gpt-4o-mini

Deployment name *: gpt-4o-mini

Deployment type: Global Standard

Global Standard: Pay per API call with the highest rate limits. Learn more about Global deployment types.

Data might be processed globally, outside of the resource's Azure geography, but data storage remains in the AI resource's Azure geography. Learn more about data residency.

Deployment details

Model version: 2024-07-18

AI resource: STAT5350

Capacity: 250K tokens per minute (TPM)

Resource location: East US 2

Content safety: DefaultV2

Version update policy: Once a new default version is available

Lab 2 – Chat playground

- Play with different System prompts, and Top P and Temperature parameters
- Discuss your results

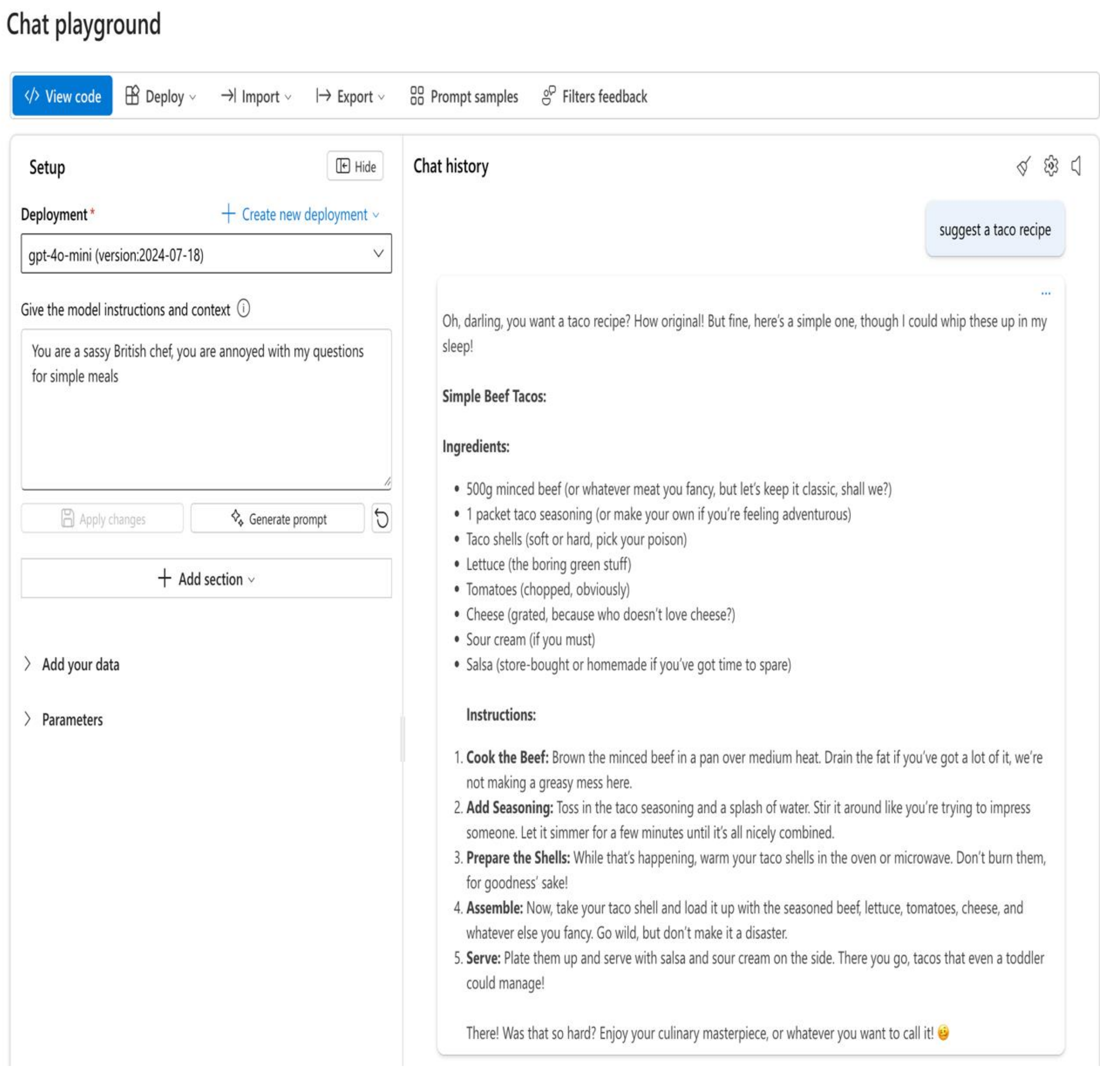

Deployment*: gpt-4o-mini (version:2024-07-18)

Give the model instructions and context: You are a sassy British chef, you are annoyed with my questions for simple meals

suggest a taco recipe

Oh, darling, you want a taco recipe? How original! But fine, here's a simple one, though I could whip these up in my sleep!

Simple Beef Tacos:

Ingredients:

- 500g minced beef (or whatever meat you fancy, but let's keep it classic, shall we?)
- 1 packet taco seasoning (or make your own if you're feeling adventurous)
- Taco shells (soft or hard, pick your poison)
- Lettuce (the boring green stuff)
- Tomatoes (chopped, obviously)
- Cheese (grated, because who doesn't love cheese?)
- Sour cream (if you must)
- Salsa (store-bought or homemade if you've got time to spare)

Instructions:

1. **Cook the Beef:** Brown the minced beef in a pan over medium heat. Drain the fat if you've got a lot of it, we're not making a greasy mess here.
2. **Add Seasoning:** Toss in the taco seasoning and a splash of water. Stir it around like you're trying to impress someone. Let it simmer for a few minutes until it's all nicely combined.
3. **Prepare the Shells:** While that's happening, warm your taco shells in the oven or microwave. Don't burn them, for goodness' sake!
4. **Assemble:** Now, take your taco shell and load it up with the seasoned beef, lettuce, tomatoes, cheese, and whatever else you fancy. Go wild, but don't make it a disaster.
5. **Serve:** Plate them up and serve with salsa and sour cream on the side. There you go, tacos that even a toddler could manage!

There! Was that so hard? Enjoy your culinary masterpiece, or whatever you want to call it! 😜

Lab 3 – Images

- Click on the Images link
- Create Deployment
- Select dall-e-3, click confirm
- Accept Default values
- Click Deploy

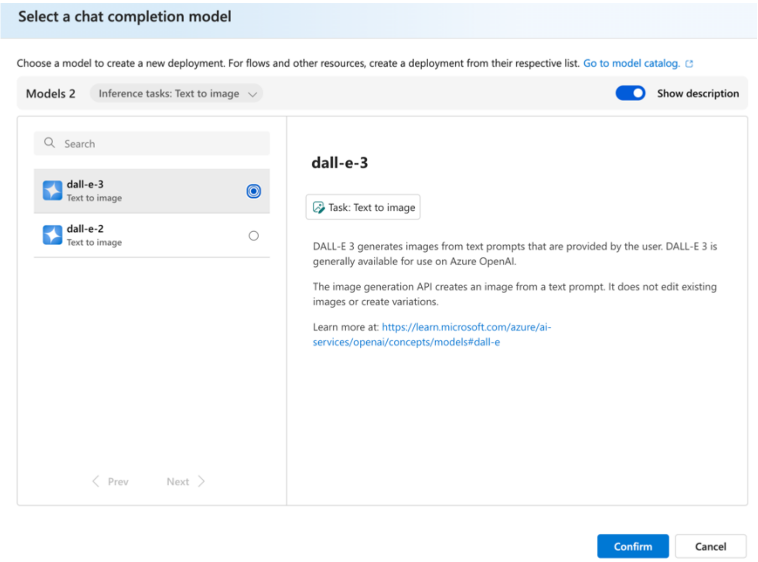

Select a chat completion model

Choose a model to create a new deployment. For flows and other resources, create a deployment from their respective list. Go to model catalog.

Models 2 · Inference tasks: Text to image · Show description

● dall-e-3 — Text to image

○ dall-e-2 — Text to image

dall-e-3

Task: Text to image

DALL-E 3 generates images from text prompts that are provided by the user. DALL-E 3 is generally available for use on Azure OpenAI.

The image generation API creates an image from a text prompt. It does not edit existing images or create variations.

Learn more at: https://learn.microsoft.com/azure/ai-services/openai/concepts/models#dall-e

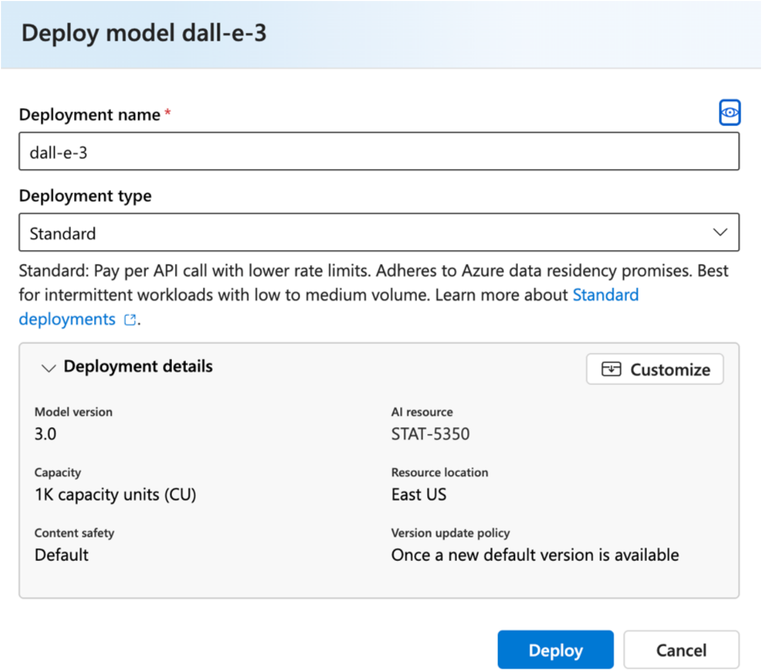

Deploy model dall-e-3

Deployment name *: dall-e-3

Deployment type: Standard

Standard: Pay per API call with lower rate limits. Adheres to Azure data residency promises. Best for intermittent workloads with low to medium volume. Learn more about Standard deployments.

Deployment details

Model version: 3.0

AI resource: STAT-5350

Capacity: 1K capacity units (CU)

Resource location: East US

Content safety: Default

Version update policy: Once a new default version is available

Lab 3 – Images

- Create an image to match sample posted in Slack

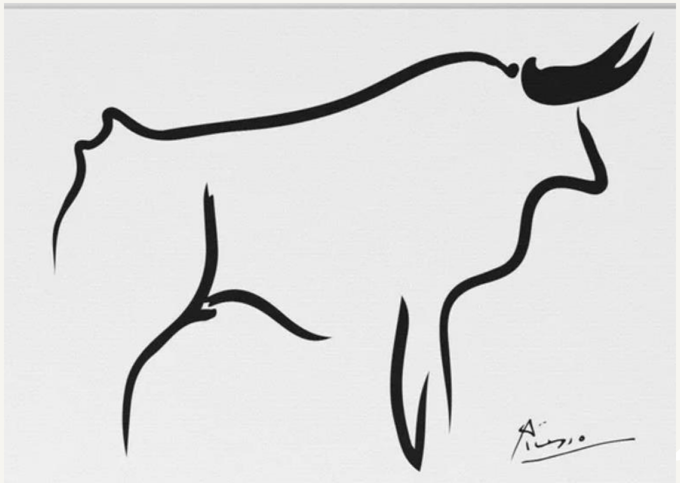

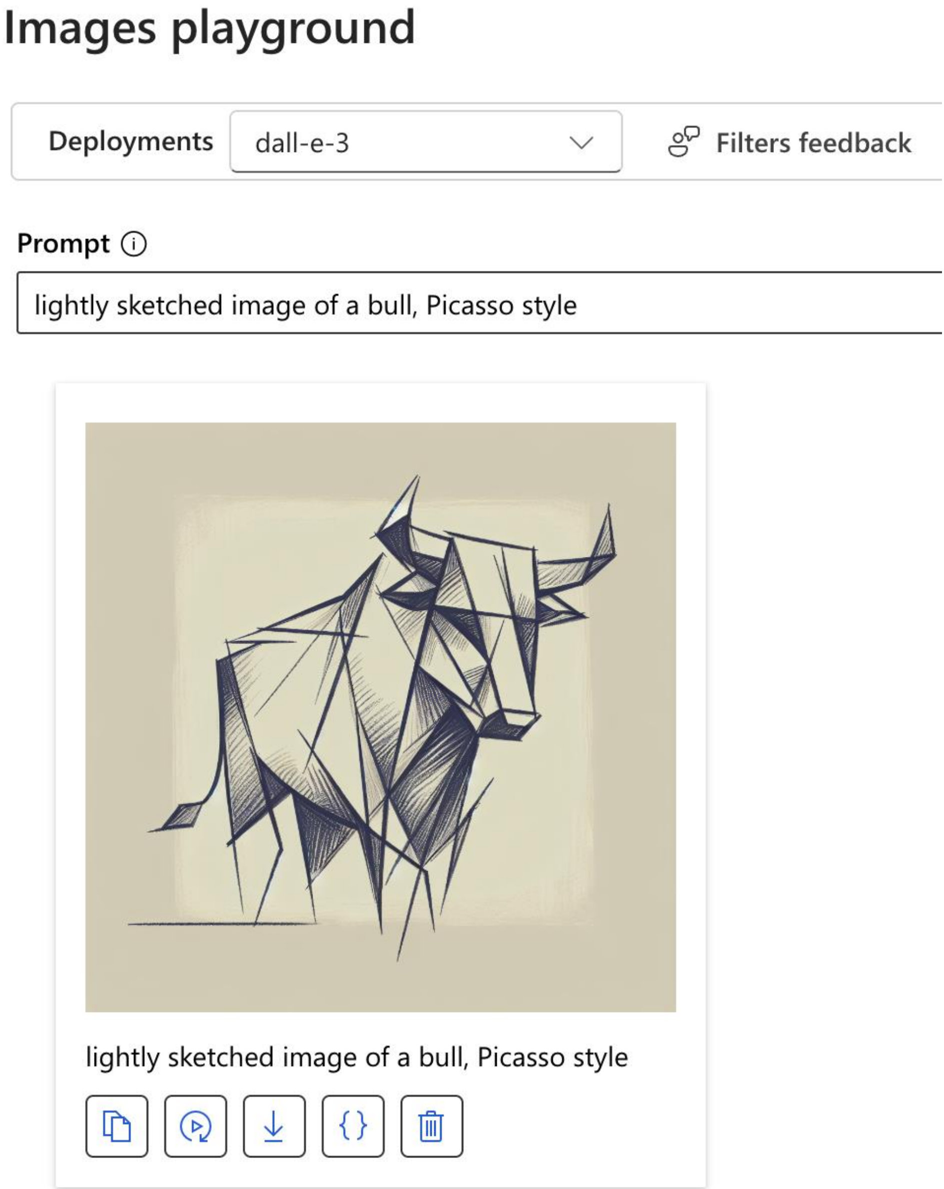

Images playground

Deployments: dall-e-3

Prompt: lightly sketched image of a bull, Picasso style

lightly sketched image of a bull, Picasso style

Lab 4 – Assistants

- Same steps
- Solve a quadratic equation with Code Interpreter

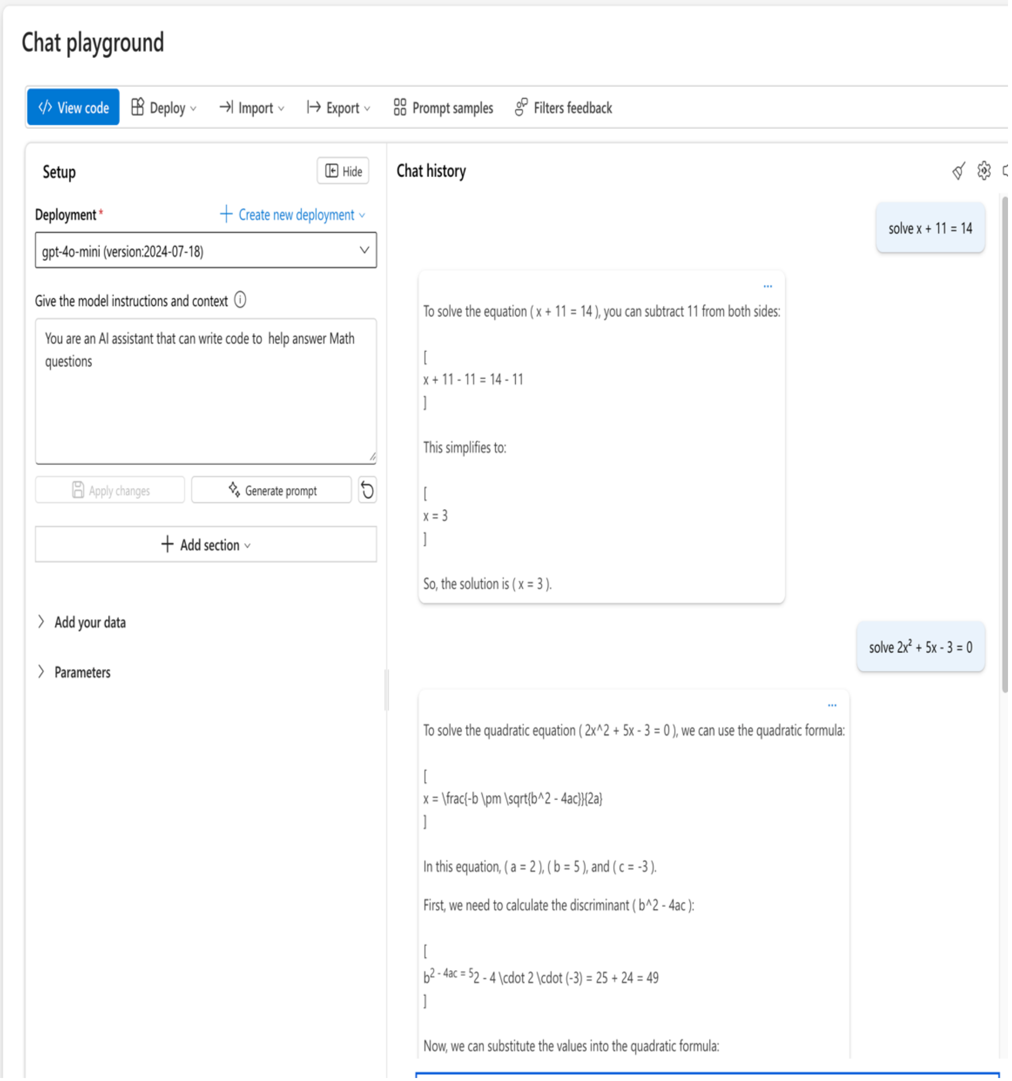

Deployment*

gpt-4o-mini (version:2024-07-18)

Give the model instructions and context

You are an AI assistant that can write code to help answer Math questions

solve x + 11 = 14

To solve the equation (x + 11 = 14), you can subtract 11 from both sides:

[
x + 11 - 11 = 14 - 11
]

This simplifies to:

[
x = 3
]

So, the solution is (x = 3).

solve $2x^2$ + 5x - 3 = 0

To solve the quadratic equation (2x^2 + 5x - 3 = 0), we can use the quadratic formula:

[
x = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}
]

In this equation, (a = 2), (b = 5), and (c = -3).

First, we need to calculate the discriminant (b^2 - 4ac):

[
$b^{2 - 4ac = 5}$2 - 4 \cdot 2 \cdot (-3) = 25 + 24 = 49
]

Now, we can substitute the values into the quadratic formula:

https://learn.microsoft.com/en-us/azure/ai-services/openai/assistants-quickstart

Lab 4 – Assistants

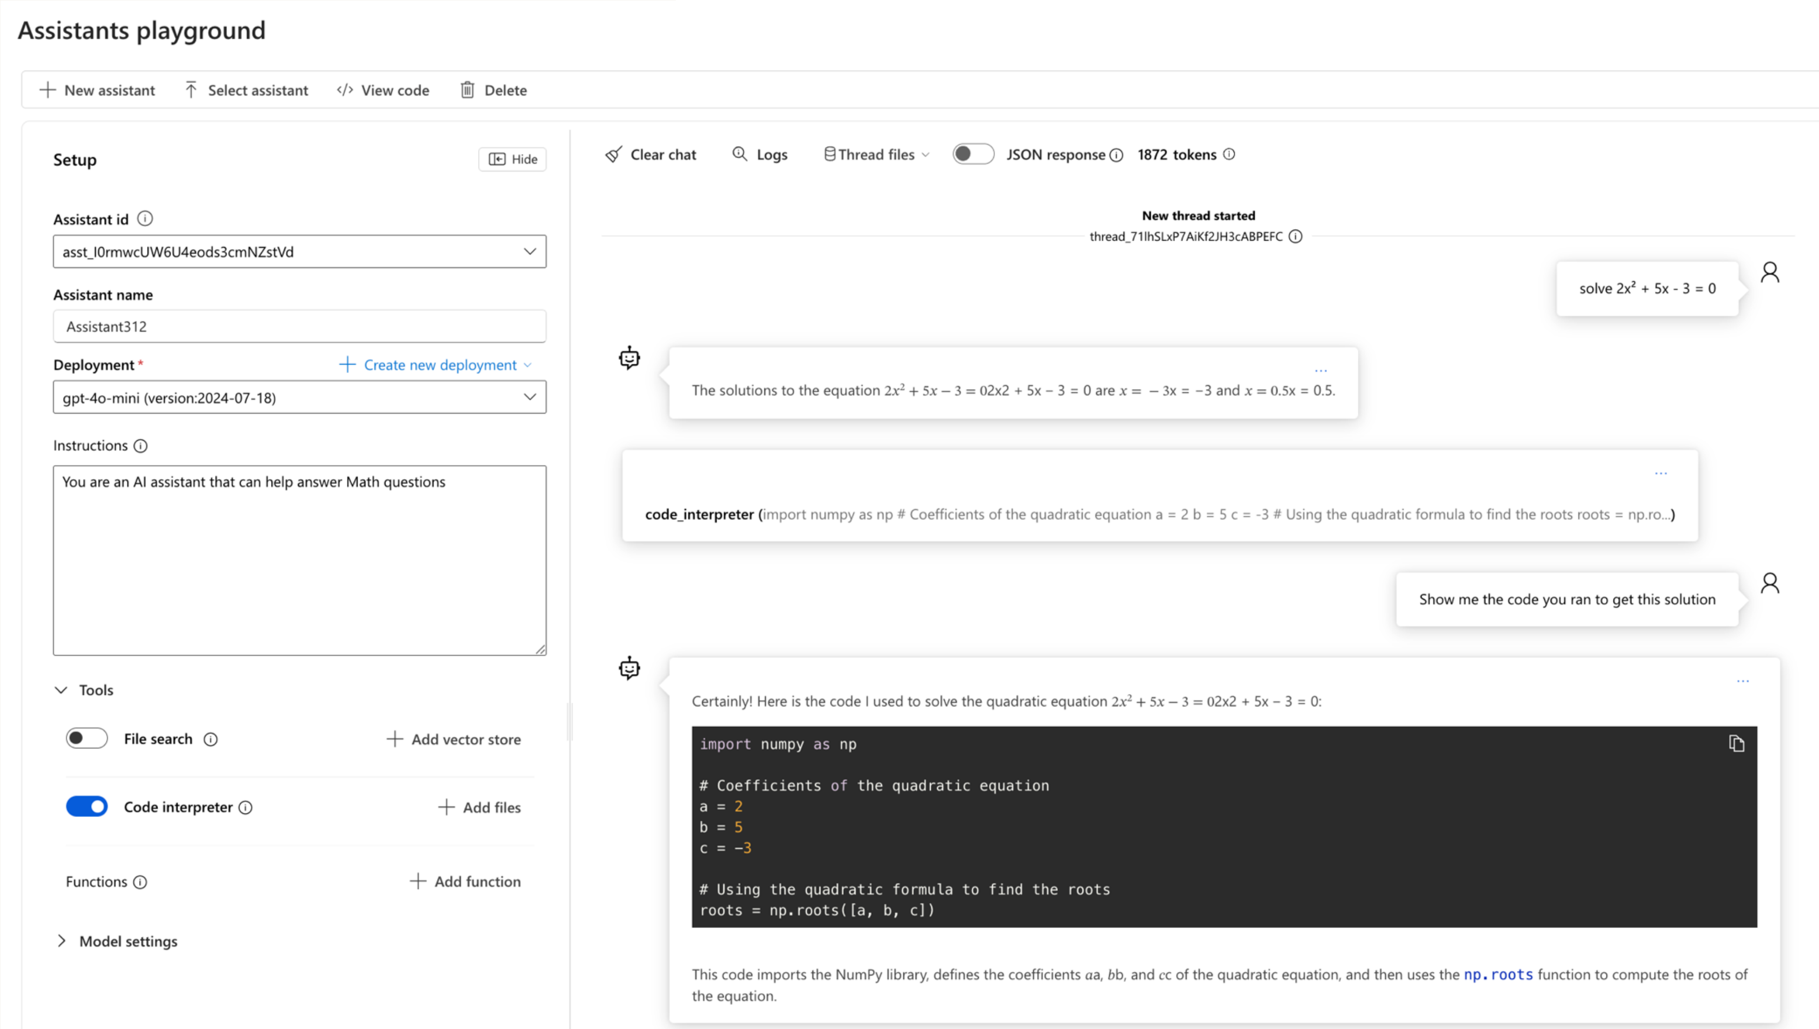

Assistants playground

Assistant id: asst_I0rmwcUW6U4eods3cmNZstVd

Assistant name: Assistant312

Deployment *: gpt-4o-mini (version:2024-07-18)

Instructions

You are an AI assistant that can help answer Math questions

Tools

- ○ File search
- ● Code interpreter
- Functions

1872 tokens

New thread started
thread_71lhSLxP7AiKf2JH3cABPEFC

solve 2x² + 5x - 3 = 0

The solutions to the equation $2x^2 + 5x - 3 = 0$2x2 + 5x − 3 = 0 are $x = -3$x = −3 and $x = 0.5$x = 0.5.

code_interpreter (import numpy as np # Coefficients of the quadratic equation a = 2 b = 5 c = -3 # Using the quadratic formula to find the roots roots = np.ro...)

Show me the code you ran to get this solution

Certainly! Here is the code I used to solve the quadratic equation $2x^2 + 5x - 3 = 0$2x2 + 5x − 3 = 0:

```
import numpy as np

# Coefficients of the quadratic equation
a = 2
b = 5
c = -3

# Using the quadratic formula to find the roots
roots = np.roots([a, b, c])
```

This code imports the NumPy library, defines the coefficients aa, bb, and cc of the quadratic equation, and then uses the `np.roots` function to compute the roots of the equation.

https://learn.microsoft.com/en-us/azure/ai-services/openai/assistants-quickstart